

LING/COMP 445, LING 645

Problem Set 3

Due before 10:05 AM on Thursday, October 19, 2021

There are several types of questions below.

- For questions involving answers in English or mathematics or a combination of the two, put your answers to the question in an **Answer** section like in the example below. You can find more information about L^AT_EX here <https://www.latex-project.org/>.
- For programming questions, please put your answers into a file called `ps3-lastname-firstname.clj`. Be careful to follow the instructions exactly and be sure that all of your function definitions use the precise names, number of inputs and input types, and output types as requested in each question.

For the code portion of the assignment, **it is crucial to submit a standalone file that runs**. Before you submit `ps3-lastname-firstname.clj`, make sure that your code executes correctly without any errors when run at the command line by typing `clojure ps3-lastname-firstname.clj` at a terminal prompt. We cannot grade any code that does not run correctly as a standalone file, and if the preceding command produces an error, the code portion of the assignment will receive a 0.

To do the computational problems, we recommend that you install Clojure on your local machine and write and debug the answers to each problem in a local copy of `ps3-lastname-firstname.clj`. You can find information about installing and using Clojure here <https://clojure.org/>.

A template Clojure file will be provided with the helper functions described below.

Once you have entered your answers, please compile your copy of this L^AT_EX into a PDF and submit

- (i) the compiled PDF renamed to `ps3-lastname-firstname.pdf`
- (ii) the raw L^AT_EX file renamed to `ps3-lastname-firstname.tex` and
- (iii) your `ps3-lastname-firstname.clj`

to the Problem Set 3 folder under ‘Assignments’ on MyCourses.

Example Problem: This is an example question using some fake math like this $L = \sum_0^\infty \mathcal{G}\delta_x$.

Example Answer: Put your answer right under the question like this $L = \sum_0^\infty \mathcal{G}\delta_x$.

Problem 1: In these exercises, we are going to be processing some natural linguistic data, the first paragraph of Moby Dick. We will first write some procedures that help us to manipulate this corpus. We will then start analyzing this data using some probabilistic models.

We will start by defining the variable `moby-word-tokens`, which is the tokens in our corpus as a Clojure list.¹

```
(def moby-word-tokens '(CALL me Ishmael . Some years ago never mind
  how long precisely having little or no money in my purse , and
  nothing particular to interest me on shore , I thought I would
  sail about a little and see the watery part of the world . It is
  a way I have of driving off the spleen , and regulating the
  circulation . Whenever I find myself growing grim about the mouth
  whenever it is a damp , drizzly November in my soul whenever I
  find myself involuntarily pausing before coffin warehouses , and
  bringing up the rear of every funeral I meet and especially
  whenever my hypos get such an upper hand of me , that it requires
  a strong moral principle to prevent me from deliberately stepping
  into the street , and methodically knocking people's hats off
  then , I account it high time to get to sea as soon as I can .
  This is my substitute for pistol and ball . With a philosophical
  flourish Cato throws himself upon his sword I quietly take to the
  ship . There is nothing surprising in this . If they but knew it
  , almost all men in their degree , some time or other , cherish
  very nearly the same feelings toward the ocean with me .))
```

Below (and in the provided template Clojure file) we have defined the function `member-of-list?`, which takes two arguments, `w` and `l`. The function returns `true` if `w` is a member of the list `l`, and `false` otherwise. For example, if `w` is `'the` and `l` is `'(the whale is)`, then the function will return `true`. In contrast, if `w` is `'the` and `l` is `'(whale is)`, then the function will return `false`.

```
(defn member-of-list? [w l]
  (if (empty? l)
    false
    (if (= w (first l))
      true
      (member-of-list? w (rest l)))))
```

Below we have defined the skeleton for the function `get-vocabulary`, which you must implement. This function takes two arguments, `word-tokens` and `vocab`. `word-tokens` is a list of words, and the function should return the list of unique words occurring in `word-tokens`. This list of unique words represents the set called the vocabulary. For example, if `word-tokens` is `'(the whale is whale the)`, then `get-vocabulary` should return `'(the whale is)`.

Fill in the missing parts of this function. When you call `(get-vocabulary moby-word-tokens '())`, you will get back a list of all of the unique words occurring in `moby-word-tokens`. Give this the name `moby-vocab`.

```
;;(defn get-vocabulary [word-tokens vocab]
;;  (if (empty? word-tokens)
;;    vocab
;;    (if (member-of-list? ;;finish this line
;;      (get-vocabulary ;;finish this line
;;      (get-vocabulary ;;finish this line
```

¹Note that commas are treated as whitespace in Clojure so they are ignored, but we left them in the list for readability.

Answer 1: Please put your answer in `ps3-lastname-firstname.clj`.

Problem 2: Define a function `get-count-of-word`. This function should take three arguments, `w`, `word-tokens`, and `count`, where `w` is a word, `word-tokens` is a list of words, and `count` is a number.

```
;;(defn get-count-of-word [w word-tokens count]
  ;;fill this in
```

When you call `(get-count-of-word w word-tokens 0)`, the function should return the number of occurrences of the word `w` in the list `word-tokens`. For example

- `(get-count-of-word 'the (list 'the 'the 'whale) 0)` should return 2.
- `(get-count-of-word 'the (list 'the 'whale) 0)` should return 1.

Write `get-count-of-word` as a recursive function. You can use the `count` argument to accumulate the words counted so far.

Answer 2: Please put the answer in `ps3-lastname-firstname.clj`.

Problem 3: Below we have defined the function `get-word-counts`, which takes two arguments, `vocab` and `word-tokens`. `vocab` is assumed to be a list of the unique words that occur in the list `word-tokens`. The function returns the number of times each word in `vocab` occurs in `word-tokens`. For example, suppose `vocab` is `'(whale the is)`, and `word-tokens` is `'(the is whale is)`. Then the function would return `(list 1 1 2)`, corresponding to the number of times `'whale`, `'the`, and `'is` occur in `word-tokens`.

```
(defn get-word-counts [vocab word-tokens]
  (let [count-word
        (fn [w] (get-count-of-word w word-tokens 0))]
    (map count-word vocab)))
```

Use the function `get-word-counts`, and the other variables we have defined, to define a variable `moby-word-frequencies`. This variable should contain the number of times each word in `moby-vocab` occurs in `moby-word-tokens`.

Answer 3: Please put your answer in `ps3-lastname-firstname.clj`.

Problem 4: In class we defined the functions `normalize`, `flip`, and `sample-categorical`. These functions are very useful for us, and are included below as well as in the Clojure template file.

```
(defn flip [p]
  (if (< (rand 1) p)
    true
    false))
```

```
(defn normalize [params]
  (let [sum (apply + params)]
    (map (fn [x] (/ x sum)) params)))
```

```
(defn sample-categorical [outcomes params]
  (if (flip (first params))
    (first outcomes)
    (sample-categorical (rest outcomes) (normalize (rest params)))))
```

Let's define a function that returns a particular probability distribution, the *uniform distribution*. The uniform distribution is the distribution which assigns equal probability to every possible outcome.

The function `create-uniform-distribution` takes a single argument, `outcomes`, which is a list of length n . The function returns a list containing the number $1/n$ repeated n times. For example, if `outcomes` is `'(the a every)`, then this function will return `'(1/3 1/3 1/3)`. This list can be interpreted as a probability distribution over the outcomes, which assigns equal probability to each of them.

```
(defn create-uniform-distribution [outcomes]
  (let [num-outcomes (count outcomes)]
    (map (fn [x] (/ 1 num-outcomes))
         outcomes)))
```

Using functions `create-uniform-distribution` and `sample-categorical`, write a function `sample-uniform-BOW-sentence` that takes a number `n` and a list `vocab`, and returns a sentence of length `n`. Each word in the sentence should be generated independently from the uniform distribution over `vocab`. For example, if `n` is 4 and `vocab` is `'(the a every)`, a possible return value for this function is `'(a the the a)`.

Note that this is a bag of words model, as defined in class. That is, we assume every element of the list is generated independently. We will call this the uniform bag of words model.

Answer 4: Please put your answer in `ps3-lastname-firstname.clj`.

Problem 5: Define a function `compute-uniform-BOW-prob`, which takes two arguments, `vocab` and `sentence`. `vocab` is the list of all words in the vocabulary, and `sentence` is a list of observed words. The function should return the probability of the sentence according to the uniform bag of words model.

For example, if `vocab` is `'(the a every)`, and `sentence` is `'(every every)`, then the function should return the number $\frac{1}{9}$.

Answer 5: Please put your answer in `ps3-lastname-firstname.clj`.

Problem 6: Using `sample-uniform-BOW-sentence` and `moby-vocab`, sample a 3-word sentence from the vocabulary of our Moby Dick corpus. This will be a sample from the uniform bag of words model for this vocabulary. Repeat this process a handful of times. For each of these 3-word sentences, use `compute-uniform-BOW-prob` to compute the probability of the sentence according to the uniform bag of words model. Are the different sentences you sampled assigned different probabilities under this model? Why is this to be expected?

Answer 6: We expect the probability of any 3-word sentence for which every word is in the `moby-vocab` to have the same probability of occurring. This is because under the bag of words approach, we assume that each word is generated independently from some distribution without reference to the other words. In this case, the distribution is uniform which means that all the words in the vocabulary are equally likely to be generated every time a word is sampled from the distribution to construct a sentence.

Problem 7: In class we looked at a more general version of the bag of words model, in which different words in the vocabulary can be assigned different probabilities. We defined a function `sample-BOW-sentence`, which returns a sentence sampled from the bag of words model that we have specified. Below we have included a slight variant of the function which we defined in class. Previously the variables `vocabulary` and `probabilities` were defined outside of the function. In the current version, they are passed in as arguments. The function is identical otherwise.

```
(defn sample-BOW-sentence [len vocabulary probabilities]
  (if (= len 0)
      '()
      ...))
```

```
(cons (sample-categorical vocabulary probabilities)
      (sample-BOW-sentence (- len 1) vocabulary probabilities))))
```

The function `sample-BOW-sentence` allows us to sample a sentence given arbitrary probabilities for the words in our vocabulary. Let's make use of this power and define a distribution over the vocabulary which is better than the uniform distribution. We will use the word frequencies for our Moby Dick corpus to *estimate* a better distribution.

Above we defined the variable `moby-word-frequencies`, which contains the frequency of every word that occurs in our Moby Dick corpus. Using `normalize` and `moby-word-frequencies`, define a variable `moby-word-probabilities`. This variable should contain probabilities for every word in `moby-vocab`, in proportion to its frequency in the text. A word which occurs 2 times should receive twice as much probability as a word which occurs 1 time.

Answer 7: Please put your answer in `ps3-lastname-firstname.clj`.

Problem 8: Using `sample-BOW-sentence`, sample a 3-word sentence from a bag of words model, in which the probabilities are set to be those in `moby-word-probabilities`. Repeat this process a number of times, and write down the sentences that you collect through this process.

Answer 8: (have growing With), (the philosophical having), (spleen of no), (involuntarily I I), (men mouth feelings).

Problem 9: Define a function `lookup-probability`, which takes three arguments, `w`, `outcomes`, and `probs`. `probs` represents a probability distribution over the elements of `outcomes`. For example, if `outcomes` is `'(the a every)`, then `probs` may be `'(0.2 0.5 0.3)`. The first number in `probs` is the probability of the first element of `outcomes`, the second number in `probs` is the probability of the second element of `outcomes`, and so on.

`lookup-probability` should look up the probability of the element `w`. For example, if `w` is `'the`, then look-up probability should return `0.2`. If `w` is `'a`, then `lookup-probability` should return `0.5`. If `w` is not in the list of `outcomes`, the function should return that its probability is zero.

Answer 9: Please put your answer in `ps3-lastname-firstname.clj`.

Problem 10: Using `lookup-probability` and `product`, define a function `compute-BOW-prob` which takes three arguments, `sentence`, `vocabulary`, and `probabilities`. The arguments `vocabulary` and `probabilities` are used to define a bag of words model with the associated probability distribution over vocabulary words. The function should compute the probability of the sentence (which is a list of words) according to the bag of words model.

This function is a generalization of the function `compute-uniform-BOW-prob` that you defined above.

Answer 10: Please put your answer in `ps3-lastname-firstname.clj`.

Problem 11: In problem 8, you collected a number of 3-word sentences. These sentences were generated from a bag of words model in which the probabilities were set to those in `moby-word-probabilities`, which reflect the relative frequency of the words in the Moby Dick corpus. Use `compute-BOW-prob` to compute the probability of these sentences according to the bag of words model. How does your answer differ from problem 6?

Choose one of the 3-word sentences that you have generated. Can you construct a different sentence which has the same probability according to the bag of words model? When computing the probability

of a sentence under a bag of words model, what information about the sentence suffices to compute this probability?

Answer 11: The probability of the 3-word sentences is not the same for all sentences since some words have higher probability of occurring (e.g. they are more frequent). Choosing '(men mouth feelings), we get a probability of $1/9129329$. Any 3-word sentence where the probability of each word is $1/209$ will do. Consider for instance '(involuntarily pausing himself). All that is needed to compute the probability of a sentence under a bag of word model is the probability of each word since every word in a sentence is generated independently.