LING/COMP 445, LING 645
Problem Set 4

Due before 10:05 AM on Thursday, November 4, 2021

There are several types of questions below.

- For questions involving answers in English or mathematics or a combination of the two, put your answers to the question in an **Answer** section like in the example below. You can find more information about LaTeX here https://www.latex-project.org/.
- For programming questions, please put your answers into a file called `ps4-lastname-firstname.clj`. Be careful to follow the instructions exactly and be sure that all of your function definitions use the precise names, number of inputs and input types, and output types as requested in each question.

  For the code portion of the assignment, **it is crucial to submit a standalone file that runs**. Before you submit `ps4-lastname-firstname.clj`, make sure that your code executes correctly without any errors when run at the command line by typing `clojure ps4-lastname-firtname.clj` at a terminal prompt. We cannot grade any code that does not run correctly as a standalone file, and if the preceding command produces an error, the code portion of the assignment will receive a 0.

  To do the computational problems, we recommend that you install Clojure on your local machine and write and debug the answers to each problem in a local copy of `ps4-lastname-firstname.clj`. You can find information about installing and using Clojure here https://clojure.org/.

Once you have entered your answers, please compile your copy of this LaTeX into a PDF and submit

(i) the compiled PDF renamed to `ps4-lastname-firstname.pdf`
(ii) the raw LaTeX file renamed to `ps4-lastname-firstname.tex` and
(iii) your `ps4-lastname-firstname.clj`

to the Problem Set 4 folder under 'Assignments' on MyCourses.

Example Problem: This is an example question using some fake math like this $L = \sum_0^\infty \mathcal{G}\delta_x$.

Example Answer:

Put your answer in the box provided like this.
Example answer is $L = \sum_0^\infty \mathcal{G}\delta_x$.

Problem 1: In this problem set, we are going to be considering a variant of the hierarchical bag-of-words model that we looked at in class. In class, we used a Dirichlet distribution to define a prior distribution over θ, the parameter vector of the bag of words model. The Dirichlet distribution is a continuous distribution on the simplex—it assigns probability density to all the uncountably many points on the simplex.

For this problem set, we will be looking at a considerably simpler prior distribution over the parameters θ. Our distribution will be *discrete*, and in particular will only assign positive probability to a finite number of values of θ. Further, for the purposes of this problem set where we will only be scoring strings rather than generating them, we will ignore the probability of the `'stop` symbol.

The probability distribution is defined in the code below:

```
(def vocabulary '(call me ishmael))

(def theta1 (list (/ 1 2 ) (/ 1 4 ) (/ 1 4 )))
(def theta2 (list (/ 1 4 ) (/ 1 2 ) (/ 1 4 )))

(def thetas (list theta1 theta2))

(def theta-prior (list (/ 1 2) (/ 1 2)))
```

Our vocabulary in this case consists of three words. Each value of θ therefore defines a bag of words distribution over sentences containing these three words. The first value of θ (`theta1`) assigns $\frac{1}{2}$ probability to the word `'call`, $\frac{1}{4}$ to `'me`, and $\frac{1}{4}$ to `'ishmael`. The second value of θ (`theta2`) assigns $\frac{1}{2}$ probability to `'me`, and $\frac{1}{4}$ to each of the other two words. The two values of θ each have prior probability of $\frac{1}{2}$. **Assume throughout the problem set that the vocabulary and possible values of θ are fixed to these values above.**

In addition to the code above we will be using some helper functions defined in class:

```
(defn score-categorical [outcome outcomes params]
  (if (empty? params)
    (throw "no matching outcome")
    (if (= outcome (first outcomes))
      (first params)
      (score-categorical outcome (rest outcomes) (rest params)))))

(defn list-foldr [f base lst]
  (if (empty? lst)
    base
    (f (first lst)
       (list-foldr f base (rest lst)))))

(defn log2 [n]
  (/ (Math/log n) (Math/log 2)))

(defn score-BOW-sentence [sen probabilities]
  (list-foldr
   (fn [word rest-score]
     (+ (log2 (score-categorical word vocabulary probabilities))
        rest-score))
   0
   sen))

(defn score-corpus [corpus probabilities]
  (list-foldr
   (fn [sen rst]
     (+ (score-BOW-sentence sen probabilities) rst))
   0
```

```
   corpus))

(defn logsumexp [log-vals]
  (let [mx (apply max log-vals)]
    (+ mx
       (log2
        (apply +
               (map (fn [z] (Math/pow 2 z))
                    (map (fn [x] (- x mx)) log-vals)))))))
```

Recall that the function `score-corpus` is used to compute the log probability of a corpus given a particular value of the parameters θ. Also recall (from the Discrete Random Variables module) the purpose of the function `logsumexp`, which is used to compute the sum of log probabilities; you should return to the lecture notes if you don't remember what this function is doing. (Note that the version of `logsumexp` here differs slightly from the lecture notes, as it does not use the `&` notation, so it takes one argument, `log-vals`, a list of log probabilities.)

Our initial corpus will consist of two sentences:

```
(def my-corpus '((call me)
                 (call ishmael)))
```

Write a function `theta-corpus-joint`, which takes three arguments: `theta`, `corpus`, and `theta-probs`. The argument `theta` is a value of the model parameters θ, and the argument `corpus` is a list of sentences. The argument `theta-probs` is a prior probability distribution over the values of θ. The function should return the **log** of the joint probability $\Pr(C = \texttt{corpus}, \theta = \texttt{theta})$.

Use the chain-rule identity discussed in class: $\Pr(C, \theta) = \Pr(C|\theta)\Pr(\theta)$. Assume that the prior distribution $\Pr(\theta)$ is defined by the probabilities in `theta-probs`, which is a list containing the prior probability of each value of θ (that is, a list with two entries, one for the probability of `theta1` and one for `theta2`).

After defining this function, you can call `(theta-corpus-joint theta1 my-corpus theta-prior)`. This will compute log joint probability of the model parameters `theta1` and the corpus `my-corpus`.

Answer 1: Please put your answer in `ps4-lastname-firstname.clj`.

Problem 2: Write a function `compute-marginal`, which takes two arguments: `corpus` and `theta-probs`. The argument `corpus` is a list of sentences, and the argument `theta-probs` is a prior probability distribution on values of θ. The function should return the **log** of the marginal likelihood of the corpus, when the prior distribution on θ is given by theta-probs. That is, the function should return $\log[\sum_{\theta \in \Theta} \Pr(\mathbf{C} = \texttt{corpus}, \Theta = \theta)]$.

Hint: Use the `logsumexp` function defined above.

After defining `compute-marginal`, you can call `(compute-marginal my-corpus theta-prior)`. This will compute the marginal likelihood of `my-corpus` (which was defined above), given the prior distribution `theta-prior`.

Answer 2: Please put the answer in `ps4-lastname-firstname.clj`.

Problem 3: Write a function `compute-conditional-prob`, which takes three arguments: `theta`, `corpus`, and `theta-probs`. The arguments have the same interpretation as in Problems 1 and 2. The function should return the **log** of the conditional probability of the parameter value `theta`, given the corpus. Remember that the conditional probability is defined by the equation:

$$\Pr(\Theta = \theta | \mathbf{C} = \texttt{corpus}) = \frac{\Pr(\mathbf{C} = \texttt{corpus}, \Theta = \theta)}{\sum_{\theta \in \Theta} \Pr(\mathbf{C} = \texttt{corpus}, \Theta = \theta)} \quad (1)$$

[Note: don't forget that your `compute-conditional-prob` should return a **log** probability.]

Answer 3: Please put your answer in `ps4-lastname-firstname.clj`.

Problem 4: Write a function `compute-conditional-dist`, which takes two arguments: `corpus` and `theta-probs`. For every value of θ in `thetas` (i.e., `theta1` and `theta2`), it should return the log conditional probability of θ given the corpus. That is, it should return a two-element list of log conditional probabilities, one for each of the two values of θ.

Answer 4: Please put your answer in `ps4-lastname-firstname.clj`.

Problem 5: Call `(compute-conditional-dist my-corpus theta-prior)`. What do you notice about the conditional distribution over values of θ? You may want to exponentiate the values you get back, so that you can see the regular probabilities, rather than the log probabilities. Explain why the conditional distribution looks the way it does, with reference to the properties of `my-corpus`. In particular, if one value of θ has higher conditional probability than the other, say why.

Answer 5: Please put your answer in the box below.

When calling `(compute-conditional-dist my-corpus theta-prior)` we get that $\Pr(\Theta = \theta_1 | \mathbf{C} = \texttt{corpus}) > \Pr(\Theta = \theta_2 | \mathbf{C} = \texttt{corpus})$. This is explained by the fact that two of the four words in `my-corpus` are `'call` and that θ_1 assigns more probability mass to the word `'call` while θ_2 assigns a smaller probability to it. θ_1 is therefore the more likely model given the corpus.

Problem 6: When you call `compute-conditional-dist`, you get back a log probability distribution over values of θ (the conditional distribution over θ given an observed corpus). This is a probability distribution just like any other. In particular, it can be used as the prior distribution over values of θ in a hierarchical bag of words model. Given this new hierarchical BOW model, we can do all of the things that we normally do with such a model. In particular, we can compute the marginal likelihood of a corpus under this model. This marginal likelihood is called a *posterior predictive distribution*.

Below we have defined the skeleton of a function `compute-posterior-predictive`, which you must complete. It takes three arguments: `observed-corpus`, `new-corpus`, and `theta-probs`. The argument `observed-corpus` is a corpus which we have observed, and are using to compute a conditional distribution over values of θ. Given this conditional distribution over θ, we will then compute the marginal likelihood of the corpus `new-corpus`. The function `compute-posterior-predictive` should return the marginal log likelihood of the new corpus given the conditional distribution on θ.

```
(defn compute-posterior-predictive [observed-corpus new-corpus theta-probs]
  (let [conditional-dist ...
    (compute-marginal ...
```

Once you have implemented `compute-posterior-predictive`, call `(compute-posterior-predictive my-corpus my-corpus theta-prior)`. What does this quantity represent? How does its value compare to the marginal likelihood that you computed in Problem 2? Why is this to be expected?

Answer 6: Please put your code in `ps4-lastname-firstname.clj` and write the text part of the answer in the box below.

The quantity computed by calling `(compute-posterior-predictive my-corpus my-corpus theta-prior)` represents the marginal log likelihood of `my-corpus` given the conditional distribution on Θ. In other words, it represents a re-weighted belief on the likelihood of the data given an updated prior belief after having observed `my-corpus`. The conditional distribution used as prior assigns more probability mass to θ_1 since it is the more likely model given `my-corpus` which leads to a higher marginal likelihood than obtained in Problem 2.

Problem 7: In the previous problems, we have written code that will compute marginal and conditional distributions *exactly*, by enumerating over all possible values of θ. In the next problems, we will develop an alternate approach to computing these distributions. Instead of computing these distributions exactly, we will approximate them using random sampling.

The following functions were defined in class, and will be useful for us going forward:

```
(defn normalize [params]
  (let [sum (apply + params)]
    (map (fn [x] (/ x sum)) params)))

(defn flip [weight]
  (if (< (rand 1) weight)
    true
    false))

(defn sample-categorical [outcomes params]
  (if (flip (first params))
    (first outcomes)
    (sample-categorical (rest outcomes)
                        (normalize (rest params)))))

(defn sample-BOW-sentence [len probabilities]
  (if (= len 0)
    '()
    (cons (sample-categorical vocabulary probabilities)
          (sample-BOW-sentence (- len 1) probabilities))))
```

Recall that the function `sample-BOW-sentence` samples a sentence from the bag of words model of length `len`, given the parameter vector `probabilities`.

Define a function `sample-BOW-corpus`, which takes three arguments: `theta`, `sent-len`, and `corpus-len`. The argument `theta` is a value of the model parameters θ. The arguments `sent-len` and `corpus-len` are positive integers. The function should return a sample corpus from the bag of words model, given the model parameters `theta`. Each sentence should be of length `sent-len` and number of sentences in the corpus should be equal to `corpus-len`. For example, if `sent-len` equals 3 and `corpus-len` equals 2, then this function should return a list of 2 sentences, each consisting of 3 words.

Hint: Use `sample-BOW-sentence`. You may also want to use the built-in function `repeatedly`.

Answer 7: Please put your answer in `ps4-lastname-firstname.clj`.

Problem 8: Below we have defined the skeleton of the function `sample-theta-corpus` which you must complete. This function takes three arguments: `sent-len` `corpus-len` and `theta-probs`. It returns a list

with two elements: a value of θ sampled from the distribution defined by `theta-probs`; and a corpus sampled from the bag of words model given the sampled θ. (The number of sentences in the corpus should equal `corpus-len`, and each sentence should have `sent-len` words in it.)

We will call the return value of this function a `theta-corpus` pair.

```
(defn sample-theta-corpus [sent-len corpus-len theta-probs]
  (let [theta ...
    (list theta ...
```

Answer 8: Please put your answer in `ps4-lastname-firstname.clj`.

Problem 9: Below we have defined some useful functions for us. The function `get-theta` takes a `theta-corpus` pair, and returns the value of `theta` in it. The function `get-corpus` takes a `theta-corpus` pair, and returns its corpus value. The function `sample-thetas-corpora` samples multiple theta-corpus pairs, and returns a list of them. In particular, the number of samples it returns equals sample-size.

```
(defn get-theta [theta-corpus]
  (first theta-corpus))

(defn get-corpus [theta-corpus]
  (first (rest theta-corpus)))

(defn sample-thetas-corpora [sample-size sent-len corpus-len theta-probs]
  (repeatedly sample-size (fn [] (sample-theta-corpus sent-len corpus-len theta-probs))))
```

We are now going to estimate the marginal likelihood of a corpus by using random sampling. Here is the general approach that we are going to use. We are going to sample some number (for example 1000) of `theta-corpus` pairs. These are 1000 samples from the joint distribution defined by the hierarchical bag of words model. We are then going to throw away the values of `theta` that we sampled; this will leave us with 1000 corpora sampled from our model.

We are going to use these 1000 sampled corpora to estimate the probability of a specific target corpus. The process here is simple. We just count the number of times that our target corpus appears in the 1000 sampled corpora. The ratio of the occurrences of the target corpus to the number of total corpora gives us an estimate of the target's probability.

More formally, let us suppose that we are given a target corpus $\mathbf{t}$. We will define the indicator function $\mathbb{1}_{\mathbf{t}}$ by:

$$\mathbb{1}_{\mathbf{t}}(c) = \begin{cases} 1, & \text{if } t = c \\ 0, & \text{otherwise} \end{cases} \tag{2}$$

We will sample n corpora $c_1, \ldots, c_n$ from the hierarchical bag of words model. We will estimate the marginal likelihood of the target corpus $\mathbf{t}$ by the following formula:

$$\sum_{\theta \in \Theta} \Pr(\mathbf{C} = \mathbf{t}, \Theta = \theta) \approx \frac{1}{n} \sum_{i}^{n} \mathbb{1}_{\mathbf{t}}(c_i) \tag{3}$$

Define a procedure `estimate-corpus-marginal`, which takes five arguments: `corpus`, `sample-size`, `sent-len`, `corpus-len`, and `theta-probs`. The argument `corpus` is the target corpus whose marginal likelihood we want to estimate. `sample-size` is the number of corpora that we are going to sample from the hierarchical model (its value was 1000 in the discussion above). The arguments `corpus-len` and `sent-len` characterize the number of sentences in the corpus and the number of words in each sentence, respectively. The argument `theta-probs` is the prior probability distribution over θ for our hierarchical model.

The procedure should return an estimate of the marginal **(not log)** likelihood of the target corpus, using the formula defined in Equation 3.

Hint: Use `sample-thetas-corpora` to get a list of samples of `theta-corpus pairs`, and then use `get-corpus` to extract the corpus values from these pairs (and ignore the `theta` values).

Answer 9: Please put your answer in `ps4-lastname-firstname.clj`.

Problem 10: Call `(estimate-corpus-marginal my-corpus 50 2 2 theta-prior)` a number of times. What do you notice? Now call `(estimate-corpus-marginal my-corpus 10000 2 2 theta-prior)` a number of times. How do these results compare to the previous ones? How do these results compare to the exact marginal likelihood that you computed in Problem 2?

Answer 10: Please put your answer in the box below.

The estimated marginal likelihood tends towards the exact marginal likelihood as we increase the value of `sample-size`. In order to get a metric for how stable the estimate was given the sample size, I computed the standard deviation for 100 calls to `(estimate-corpus-marginal` with sample sizes of 50 and 10000 as illustrated by the following table.

sample-size	mean-estimate	std-dev
50	0.0134	0.0176
10000	0.0118	0.0011

We can see that the the mean estimate of the marginal likelihood for sample sizes of 10000 is closer to the exact marginal likelihood of 0.0117 computed in Problem 2 and is stable as indicated by the very small standard deviation.

Problem 11: These functions will be useful for us in the next problem.

```
(defn get-count [obs observation-list count]
  (if (empty? observation-list)
    count
    (if (= obs (first observation-list))
      (get-count obs (rest observation-list) (+ 1 count))
      (get-count obs (rest observation-list) count))))

(defn get-counts [outcomes observation-list]
  (let [count-obs
        (fn [obs] (get-count obs observation-list 0))]
    (map count-obs outcomes)))
```

In Problem 9, we introduced a way of approximating the marginal likelihood of a corpus by using random sampling. We can similarly approximate a conditional probability distribution by using random sampling.

Suppose that we have observed a corpus **c**, and we want to compute the conditional probability of a particular θ. We can approximate this conditional probability as follows. We first sample n theta-corpus pairs. We then remove all of the pairs in which the corpus does not match our observed corpus **c**. We finally count the number of times that θ occurs in the remaining theta-corpus pairs, and divide by the total number of remaining pairs. This process is an example of *rejection sampling*.

Define a function `rejection-sampler` which has the following form:

```
(defn rejection-sampler
  [theta observed-corpus sample-size sent-len corpus-len theta-probs]
  ...
)
```

This function should use the rejection sampling method (as described above) to estimate the conditional probability of `theta`, given that we have observed the corpus `observed-corpus`. The function must estimate this conditional probability by taking `sample-size` samples (you may assume this argument is a positive integer) from the joint distribution on `theta-corpus` pairs. The procedure should filter out any `theta-corpus` pairs in which the corpus does not equal the observed corpus. If there are no remaining pairs after filtering, then the function should return `nil`. Otherwise, it should then count the number of times that `theta` occurs in the remaining pairs, and divide by the total number of those pairs.
Hint: Use `get-count` or `get-counts` to count the number of occurrences of `theta`.

Answer 11: Please put your answer to the coding problem in `ps4-lastname-firstname.clj`.

Problem 12: Call `(rejection-sampler theta1 my-corpus 100 2 2 theta-prior)` a number of times. What do you notice? Try with larger sample sizes (such as $200, 500, 1000 \ldots$). How large does `sample-size` need to be until you get a stable estimate of the conditional probability of `theta1`? Why does it take so many samples to get a stable estimate?

Answer 12: Please answer the questions in the box below.

In order to determine the sample size value at which the estimate becomes stable, I computed the standard deviation of 100 calls to `rejection-sampler` for all the values in the following table:

sample-size	std-dev
100	0.4264
500	0.2391
1000	0.1543
5000	0.0618
10000	0.0461
20000	0.0261

Given the results above, it seems that we get a relatively stable estimate of the conditional probability for sample sizes that are equal to or larger than 20000. We can explain this tendency for the estimate to stabilize using the Law of Large Numbers which says that "the average of the results obtained from a large number of trials should be close to the expected value and will tend to become closer to the expected value as more trials are performed" (theorem definition taken from Law of Large Numbers). In this specific case, the sample size needs to be very large since we first reject most of the samples to compute the estimated conditional probability of `theta1` which implies that we only ever use a small fraction of the total sample to compute our estimate.